

# 1. Learning algorithm

## 1.1. Overview of the technique

The agent is implemented using the Twin-delayed deep deterministic policy gradient (TD3 for short) method. The paper for this algorithm can be found here:

[Addressing Function Approximation Error in Actor-Critic Methods \(arxiv.org\)](https://arxiv.org/abs/1802.09454)

This algorithm is an upgraded version of the Deep deterministic policy gradient (DDPG). A quick recap, DDPG is an algorithm which concurrently learns a Q-function and a policy. It uses off-policy data and the Bellman equation to learn the Q-function, and uses the Q-function to learn the policy.

DDPG is an off-policy algorithm and can only be used in continuous action space.

TD3 added these following extensions to DDPG:

- Use Clipped Double Q-learning technique. Concretely, TD3 use 2 critics network to calculate Q-value and the one with smaller value will be used to perform loss update
- Use “Delayed” Policy Update. TD3 updates the policy (and target networks) less frequently than the Q-function. The paper recommends one policy update for every two Q-function updates
- Apply Target Policy Smoothing. TD3 adds noise to the target action, to make it harder for the policy to exploit Q-function errors by smoothing out Q along changes in action

These changes resolve DDPG problem of overestimating Q-value as well as hyperparameters tuning

The algorithm in detail:

---

**Algorithm 1** Twin Delayed DDPG


---

- 1: Input: initial policy parameters  $\theta$ , Q-function parameters  $\phi_1, \phi_2$ , empty replay buffer  $\mathcal{D}$
- 2: Set target parameters equal to main parameters  $\theta_{\text{targ}} \leftarrow \theta, \phi_{\text{targ},1} \leftarrow \phi_1, \phi_{\text{targ},2} \leftarrow \phi_2$
- 3: **repeat**
- 4:   Observe state  $s$  and select action  $a = \text{clip}(\mu_\theta(s) + \epsilon, a_{\text{Low}}, a_{\text{High}})$ , where  $\epsilon \sim \mathcal{N}$
- 5:   Execute  $a$  in the environment
- 6:   Observe next state  $s'$ , reward  $r$ , and done signal  $d$  to indicate whether  $s'$  is terminal
- 7:   Store  $(s, a, r, s', d)$  in replay buffer  $\mathcal{D}$
- 8:   If  $s'$  is terminal, reset environment state.
- 9:   **if** it's time to update **then**
- 10:     **for**  $j$  in range(however many updates) **do**
- 11:       Randomly sample a batch of transitions,  $B = \{(s, a, r, s', d)\}$  from  $\mathcal{D}$
- 12:       Compute target actions
$$a'(s') = \text{clip}(\mu_{\theta_{\text{targ}}}(s') + \text{clip}(\epsilon, -c, c), a_{\text{Low}}, a_{\text{High}}), \quad \epsilon \sim \mathcal{N}(0, \sigma)$$
- 13:       Compute targets
$$y(r, s', d) = r + \gamma(1 - d) \min_{i=1,2} Q_{\phi_{\text{targ},i}}(s', a'(s'))$$
- 14:       Update Q-functions by one step of gradient descent using
$$\nabla_{\phi_i} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi_i}(s, a) - y(r, s', d))^2 \quad \text{for } i = 1, 2$$
- 15:       **if**  $j \bmod \text{policy\_delay} = 0$  **then**
- 16:         Update policy by one step of gradient ascent using
$$\nabla_{\theta} \frac{1}{|B|} \sum_{s \in B} Q_{\phi_1}(s, \mu_\theta(s))$$
- 17:         Update target networks with
$$\begin{aligned} \phi_{\text{targ},i} &\leftarrow \rho \phi_{\text{targ},i} + (1 - \rho) \phi_i \\ \theta_{\text{targ}} &\leftarrow \rho \theta_{\text{targ}} + (1 - \rho) \theta \end{aligned} \quad \text{for } i = 1, 2$$
- 18:       **end if**
- 19:     **end for**
- 20:   **end if**
- 21: **until** convergence


---

## 1.2. Network architecture

The network architecture for actor and critic are described below (the number is # neurons at each hidden layers)



Actor network structure




Critic network structure



### 1.3. Hyperparameters select


Name	Value	Meaning
seed	73	Random seed for stable result
actor_lr	0.0001	Learning rate for actor's Adam optimizer
critic_lr	0.0001	Learning rate for critic's Adam optimizer
gamma	0.99	Discount factor
tau	0.005	Soft update target network hyperparameter
buffer_size	1000000	Number of transitions stored in replay buffer
batch_size	100	Batch size for learning phase of the agent

## 2. Plot of rewards



## 3. Ideas for future works

The result I have shown is quite good, which solve the problem in around 200 episodes. To be honest, I have played around with DDPG before TD3 and the result is terrible. I cannot solve the problem with 5000 episodes. Here is the plot of DDPG:



Although the result I got from TD3 is good enough, I think some modifications will make the result even better:

- Tuning hyperparameters such as learning rate and the architecture of the actor and critic network probably enhance the learning ability of the agent
- Moreover, many other algorithms can be used in order to help the agent learn better such as Distributed Distributional DDPG (D4PG), Proximal Policy Optimisation (PPO) or Soft Actor Critic (SAC)